

Quality per unit of compute
(bubble area $\propto$ trainable parameters)

